

Report - Linear system of equations solvers

ACSE5 – Advanced Programming Assignment

Gabriel Lipkowitz, Albert Celma Ortega, Nicolas Trinephi

February 2, 2020

Report overview

- Approach
- Core concepts
- Implementation
- Method analysis: Performance
- Takeaways on strategy and software design

1. Approach

Solving complex linear systems of equations quickly and efficiently represents an increasing field of research, and its applications have rapidly grown in the last decades jointly with computational capacity improvements. We decided to first develop a set of iterative methods for solving dense, sparse and banded matrices. From there, we decided to expand into direct methods for dense matrices and finally developed the adaptation for sparse matrices (LU decomposition and Cholesky methods).

2. Core concepts

Each of our methods were built in the developed `Matrix`, `CSRMatrix`, and `BandedMatrix` classes, declaring methods in the respective header files. For testing purposes, to ensure that our solvers indeed solve an $A\mathbf{x} = \mathbf{b}$ type system of linear equations, we included a `Matrix_Tester` class and its own corresponding `.cpp` and `.h` files. A first important programming technique for our implementation was *inheritance*. As subclasses of the dense `Matrix` class, the sparse and banded matrix classes inherited all the properties of a generic dense matrix, such as public variables like columns and rows. *Polymorphism* was also of crucial importance. Namely, several functions in the dense, sparse, and banded matrix classes completed the same operation (e.g. matrix column-vector multiplication, matrix row-vector multiplication, sparse and dense matrix-matrix multiplication, along with the Jacobi, Gauss-Seidel, SOR, conjugate gradient, and Cholesky solvers). Thanks to polymorphism, when we called these functions, our programme could automatically implement the correct method corresponding to each class. Finally, all our functions, for all three of our matrix classes, were templated so that the matrices can contain integers, doubles, floats, etc.

3. Implementation

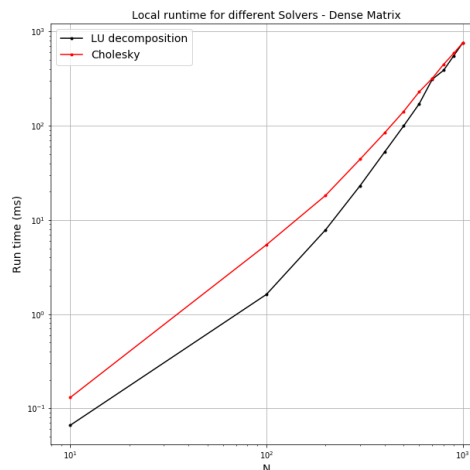
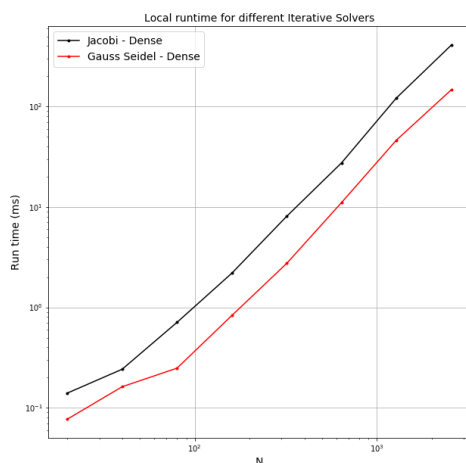
While we initially tested our solvers with simple matrices generated manually, we also wrote functions to randomly generate dense, sparse, and banded matrices, of different sizes, and sparsities and bandwidths, if applicable. Sparse matrices were stored in the old Yale format, whereas banded matrices were stored as an $n \times b$ array, where n is the number of rows and b is the bandwidth. Matrices could be filled randomly such that they were diagonally dominant (see iterative solvers' requirements for convergence, below) and/or symmetric positive definite (see conjugate gradient solver's requirements, below).

The simple iterative solvers – Jacobi, Gauss-Seidel, and Successive Over-Relaxation (SOR) – all can converge to the solution for diagonally dominant matrices. Our matrix generation functions, described above, accomplished this by enforcing the diagonal elements to be artificially large.

The conjugate gradient method requires, for convergence, that the matrix be symmetric positive definite (SPD), which is defined by the following equation:

$$\mathbf{z}^T \mathbf{A} \mathbf{z} > 0$$

To generate an SPD matrix, as a function we wrote did, we first generated a lower triangular matrix, with nonzero diagonal elements, and then multiplied it by its transpose, an upper triangular matrix also with nonzero diagonal elements. To perform these operations, we



For dense matrices, as shown below, the Cholesky method demonstrated $O(n^2)$ run time, i.e. as the matrix size increased by a factor of 10, run time increased by a factor of 100.

Description of LU Factorization:

Case 1

$$\begin{bmatrix} 5 & 3 & 0 & 0 \\ 0 & 2 & 0 & 1 \\ 1 & 0 & 4 & 0 \\ 0 & 3 & 0 & 1 \end{bmatrix} \rightarrow \begin{bmatrix} 5 & 3 & 0 & 0 \\ 0 & 2 & 0 & 1 \\ 0 & 0 & 4 & 0 \\ 0 & 3 & 0 & 1 \end{bmatrix} \rightarrow \begin{bmatrix} 5 & 3 & 0 & 0 \\ 0 & 2 & 0 & 1 \\ 0 & 0 & 4 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix} LU^{(2/2)} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0.2 & 0 & 1 & 0 \\ 0 & 1.5 & 0 & 1 \end{bmatrix} \begin{bmatrix} 5 & 3 & 0 & 0 \\ 0 & 2 & 0 & 1 \\ 0 & 0 & 4 & 0 \\ 0 & 0 & 0 & -0.5 \end{bmatrix} \text{Correct}$$

Case 2

$$\begin{bmatrix} 5 & 0 & 3 & 0 \\ 0 & 2 & 0 & 1 \\ 1 & 0 & 4 & 0 \\ 5 & 3 & 0 & 1 \end{bmatrix} \rightarrow \begin{bmatrix} 5 & 0 & 3 & 0 \\ 0 & 2 & 0 & 1 \\ 0 & 0 & 3.4 & 0 \\ 0 & 3 & -3 & 1 \end{bmatrix} \rightarrow (*) \begin{bmatrix} 5 & 3 & 0 & 0 \\ 0 & 2 & 0 & 1 \\ 0 & 0 & 4 & 0 \\ 0 & 0 & -3 & -0.5 \end{bmatrix} L = LU^{(1/2)**} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0.2 & 0 & 1 & 0 \\ 1 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 5 & 3 & 0 & 0 \\ 0 & 2 & 0 & 1 \\ 0 & 0 & 4 & 0 \\ 0 & 0 & 0 & -0.5 \end{bmatrix} \text{Partial}$$

*Appears non zero entries and this represents an update of the *sparse matrix* vectors

** second iteration not reached completely correctly

1. Areas for improvement

1. Mathematical models

We feel happy and proud of the broad range of methods that we have built. We consider that we had the right approach focusing on iterative methods and then building direct methods. With respect to direct methods, we covered LU decomposition, and Cholesky. These two are a good sample of direct methods. We feel that although significant improvement was made for LU decomposition as well as Cholesky for sparse matrix notation, we should have had a different strategy to be more efficient in managing time.

2. Effective team work

The most important takeaway from this course is definitively how to allocate work effectively to optimise our output. We decided to run on parallel for LU decomposition sparse as well as Cholesky thinking that we could work together having both solvers done and then work to get forward and backwards substitution done. But Cholesky and LU decomposition implementation were further apart than the actual concepts had allowed us to foresee. These methods for sparse matrices are incomplete but a good effort was put into trying to complete them. The take away is to work parallel only if tasks are not superimposable, or said differently, that once one both are finished and assembled, you can have a working piece of software.

References:

S. Parter, *The use of linear graphs in Gauss Elimination*, SIAM Rev., 3 (1961), pp. 119-130.

Bank, Randolph and Douglas, Craig (1993) "Sparse Matrix Multiplication Package (SMMP)" *Advances in Computational Mathematics* 1.